

JVM Performance Comparison for JDK 21

Ionut Balosin

www.IonutBalosin.com | ionut.balosin@gmail.com | [@ionutbalosin](https://twitter.com/ionutbalosin)

Agenda

01 Introduction

02 Benchmarks

03 Conclusions

04 Challenges & Lessons Learned

05 Future Work

Introduction

01

Ionut Balosin



Software Architect @  **Raiffeisen Bank International**



Technical Trainer



Security Champion



Oracle ACE Associate



Blogger



Speaker

www.IonutBalosin.com | ionut.balosin@gmail.com | [@ionutbalosin](https://twitter.com/ionutbalosin)

My Training Catalogue

Software Architecture Essentials

Java Performance Tuning

Designing High-Performance, Scalable, and Resilient Applications

Application Security for Java Developers

Training figures: 80+ sessions | 900+ trainees | 1300+ hours | 10+ clients | 4+ countries


Conference figures: 35+ sessions | 14+ countries

www.IonutBalosin.com/training

Thanks to Florin Blanaru (co-author)



 Senior Software Engineer @  **AXELERA**
ARTIFICIAL INTELLIGENCE

 TornadoVM - ex contributor

 Student of the year Award from the RISC-V Foundation, 2019

Compilers  Language Runtimes
Performance Analysis & Tuning

GitHub: @gigiblender | X: @gigiblender | Mastodon: @gigiblender

JVM Performance Comparison for JDK 21

Context

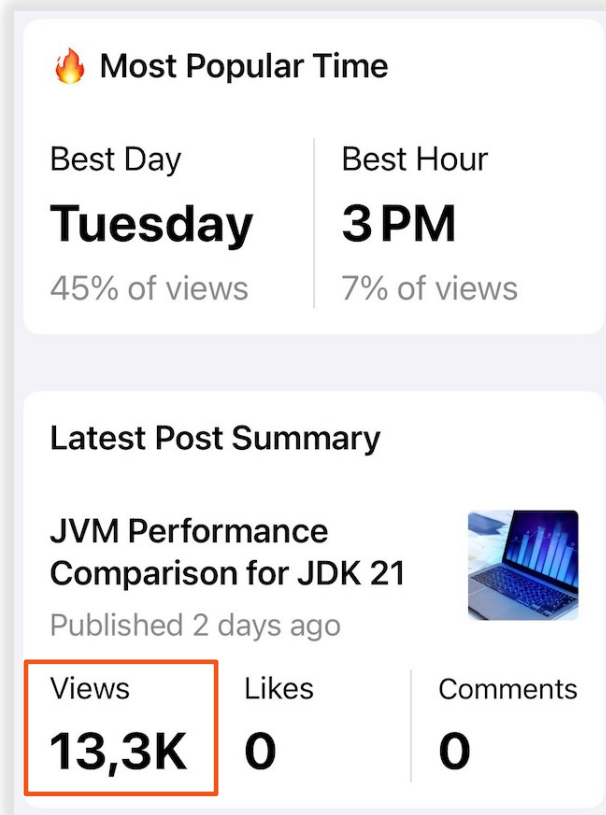
The current article describes a series of Java Virtual Machine (JVM) benchmarks with a primary focus on top-tier Just-In-Time (JIT) compilers, such as C2 JIT, and Graal JIT. The benchmarks are structured in three distinct (artificial) categories:

1. **Compiler:** This category is dedicated to assessing JIT compiler optimizations by following specific handwritten code patterns. It assesses common optimizations found in compilers, including inlining, loop unrolling, escape analysis, devirtualization, null-check elimination, range-check elimination, dead code elimination, vectorization, etc.
2. **Api:** This category includes benchmarks targeting common APIs from both the Java Platform, Standard Edition (Java SE) (e.g., `java.io`, `java.nio`, `java.net`, `java.security`, `java.util`, `java.text`, `java.time`, etc.) and the Java Development Kit (JDK) (e.g., `jdk.incubator.vector`, etc.).
3. **Miscellaneous:** This category covers a broader spectrum of classical programs (e.g., Dijkstra's shortest path, factorial, Fibonacci, Game of Life, image rotation, knapsack problem, N queens, palindrome, Huffman coding/encoding, Lempel-Ziv-Welch compression, etc.) using different techniques (e.g., dynamic programming, greedy algorithms, backtracking, divide and conquer, etc.), various programming styles (e.g., iterative, functional), and high-level Java APIs (e.g., streams, lambdas, fork-join, collections, etc.).

The categorization is for informative purposes to better organize and direct the focus of our benchmarks, ranging from low-level (compiler benchmarks) to high-level (API and Miscellaneous) benchmarks.

For this report we aggregated in total a number of **1112 benchmark runs**, including all three categories.

The traffic during the first two days after we launched the article exploded 🚀



JMVs / JIT Compilers from JDK 21



vs.



vs.




OpenJDK
C2 JIT

Oracle GraalVM
Graal JIT

GraalVM CE
Graal JIT


Hardware/OS Configuration



x86_64

Dell XPS 15 7590

CPU	Intel Core i7-9750H 6-Core
MEMORY	32GB RAM
OS / kernel	Ubuntu 20.04 LTS / 5.15.0-101-generic
	disabled hyper-threading and turbo-boost



arm64

Apple MacBook Pro

CPU	M1 Chip 10-Core, 16-Core Neural Engine
MEMORY	32GB RAM
OS / kernel	macOS Ventura 13.6.1 / Darwin kernel version 23.4.0

Software Configuration



1112 benchmark runs (in the entire JDK suite)



JMH v1.37 (5x10s warm-up it., 5x10s measurement it., 5 JVM forks **[*]**)



OpenJDK 21 | GraalVM CE 21+35.1 | Oracle GraalVM 21+35.1



Total execution time (per JDK suite/per JVM): \approx 6.5 days

[*] – each JVM fork starts with “*-Xms4g -Xmx4g -XX:+AlwaysPreTouch*”

Benchmarks

02

Infrastructure Baseline Benchmark

Used as a baseline to assess the infrastructure overheads.
Should be the same between the JVMs for a fair comparison.

References: [[article](#)][[code source](#)]



Dead Argument Elimination Benchmark

Assesses the removal of arguments that are directly dead, as well as arguments that are passed into function calls as dead arguments of other functions.

References: [[article](#)][[code source](#)]



Dead Local Allocation Store Benchmark

Checks how the compiler handles dead allocations.

Dead allocation is an allocation that is not used by subsequent instructions.

References: [[article](#)][[code source](#)]



Lock Coarsening Benchmark

Tests how the compiler can effectively coarsen/merge adjacent locks.
Optimization useful to reduce the overhead of object locking/unlocking.

References: [[article](#)][[code source](#)]



NPE Throw Benchmark

Tests the implicit vs explicit throw and catch of NPE in a hot loop.
The callee is never inlined into the caller.

References: [[article](#)][[code source](#)]



Scalar Replacement Benchmark

Tests the ability of the compiler for perform escape analysis and scalar replacement.

References: [[article](#)][[code source](#)]



String Concatenation Benchmark

Checks measuring the performance of various concatenation methods using different data types (e.g., String, int, float, char, long, double, boolean, Object):
StringBuilder, *StringBuffer*, *String.concat()*, *plus operator*, *StringTemplate*

References: [[article](#)][[code source](#)]



Enum Values Lookup Benchmark

Iterates through the *enum* values and returns the value that matches a lookup value. This pattern is commonly seen in business applications, where microservices with RESTful APIs defined in *OpenAPI/Swagger* use *enums*.

References: [[article](#)][[code source](#)]



Conclusions

03


Geometric Mean

$$\left(\prod_{i=1}^n x_i \right)^{\frac{1}{n}} = \sqrt[n]{x_1 x_2 \dots x_n}$$

“How to not lie with statistics: the correct way to summarize benchmark results”

– Philip J Fleming, John J Wallace

JIT Geometric Mean




x86_64

Note: This is purely informative to have a high-level understanding of the overall benchmark scores.

Rank	JIT	Normalized Geometric Mean	Number of Benchmarks	Unit
1	Oracle GraalVM	0.77	1112	ns/op
2	C2	1	1112	ns/op
3	GraalVM CE	1.03	1112	ns/op

The first in the row is the fastest, and the last in the row is the slowest.

JIT Geometric Mean



arm64

Note: This is purely informative to have a high-level understanding of the overall benchmark scores.

Rank	JIT	Normalized Geometric Mean	Number of Benchmarks	Unit
1	Oracle GraalVM	0.83	1112	ns/op
2	C2	1	1112	ns/op
3	GraalVM CE	1.08	1112	ns/op

The first in the row is the fastest, and the last in the row is the slowest.

Oracle GraalVM JIT



Key Strengths

- ✓ improved partial escape analysis
- ✓ more aggressive inlining heuristics (including polymorphic inlining, recursive method inlining, constructor inlining, etc.)
- ✓ more compact TLAB allocation code for grouped allocations
- ✓ extended vectorization support
- ✓ loop unrolling
- ✓ able to devirtualize and inline call sites up to four different targets
- ✓ (in general) cleaner and more compact/efficient CPU assembly instructions



Key Strengths

- ✓ vectorization support
- ✓ extended intrinsics support
- ✓ efficient exceptions handling mechanism when the same implicit exception is thrown multiple times in the hot path
- ✓ loop unrolling
- ✓ able to devirtualize and inline call sites up to two different targets

GraalVM CE JIT



Key Strengths

- ✓ improved partial escape analysis
- ✓ (in general) better inlining heuristics than C2 JIT
- ✓ able to devirtualize and inline call sites with a higher number of targets(, even exceeding Oracle GraalVM JIT)




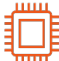

Challenges & Lessons Learned

04

- ① **Microbenchmarking is not about numbers**, without a proper understanding of what happens, the benchmark has no value.
- ② Microbenchmarking is **not always a good predictor** for large scale applications.
- ③ There are **differences between different architectures** (e.g., x86_64, arm64).
- ④ **Microbenchmarking for GC is (in general) misleading**. We initially had a few GC-targeted benchmarks, but we dropped them.

Future Work

05

-  Keep on enhancing the completeness of the benchmark suite.
-  Release the report for the next JDK LTS release(s). We have already successfully done so for JDK 17 and JDK 21.
-  Try to include the Falcon JIT Compiler (from Azul Prime VM).
-  Potentially extend it to RISC-V architecture, besides x86_64 and arm64.
-  **If you want to contribute to this project, please feel free to reach out to us.**

Thank You

Additional Resources

Code Source: [<https://github.com/ionutbalosin/jvm-performance-benchmarks>]

JDK 21 article: [<https://ionutbalosin.com/2024/02/jvm-performance-comparison-for-jdk-21>]

JDK 17 article : [<https://ionutbalosin.com/2023/03/jvm-performance-comparison-for-jdk-17>]

Skill Up Now: [<https://ionutbalosin.com/training>]